THESIS — Investment Decision Intelligence Platform

Product Ideation Document v1.0

THE ONE-LINER

Thesis is the operating system for how India invests — not what to buy, but whether your reasons for buying are any good.

TABLE OF CONTENTS

1. The Core Insight
2. Market Context & Why Now
3. What Thesis Actually Is
4. The Five Engines (Product Deep-Dive)
5. Cross-Asset Unified View
6. Data Architecture & Sources
7. SEBI Compliance Model
8. Monetization & Unit Economics
9. Competitive Landscape & Moat
10. Go-To-Market Strategy
11. Technical Architecture
12. MVP Scope & Phased Roadmap
13. Risk Register
14. Team Composition
15. Appendix: Indian Market Data Sources

---

1. THE CORE INSIGHT

Every investing tool in India answers the same question: **"What does the data say about this stock?"**

Screener.in gives you financial ratios. Tickertape gives you screeners and scorecards. Smallcase gives you curated baskets. Zerodha's Varsity teaches you what PE ratio means. Sensibull helps you build option strategies.

None of them answer the harder question:

> **"Is my reasoning for this investment actually sound — and will I know when it stops being sound?"**

This is the gap. Indian retail investors (190M+ demat accounts as of 2026) don't lose money because they lack data. They lose money because of:

- **Thesis drift —** They bought HDFC Bank for its loan book quality, but they're now holding it because "it'll come back" after the thesis already broke.
- **Invisible concentration —** They own 4 mutual funds and 8 direct stocks and think they're diversified, but 47% of their actual capital exposure is in 5 companies.
- **No exit framework —** They know exactly why they bought. They have zero framework for when to sell.
- **IPO FOMO —** They apply to every IPO because of GMP hype on Telegram, without reading a single line of the DRHP or comparing it to already-listed peers.
- **Recency bias —** They overweight recent winners and ignore base rates and cyclical patterns.

Thesis doesn't tell you what to buy. It tells you whether your *own* investment thesis is rigorous, whether your portfolio is actually doing what you think it's doing, and exactly when your reasons for holding something stop being valid.

The mental model: **Thesis is Grammarly for investment decisions.** Grammarly doesn't write for you — it makes your writing better by catching errors you can't see yourself. Thesis does the same for your investment process.

---

2. MARKET CONTEXT & WHY NOW

The Demand Side

- **190M+ demat accounts** in India (CDSL + NSDL), up from 40M in 2020
- **₹25,000 Cr+ monthly SIP flows —** retail money pouring in on autopilot
- **Average Indian retail investor holds 3-7 mutual funds + 5-15 direct stocks —** creating massive hidden overlap

- **IPO frenzy**: 2024-2026 saw 200+ mainboard IPOs, with retail oversubscription averaging 10-40x. Most retail investors apply based on GMP (Grey Market Premium) signals from Telegram groups, not fundamental analysis
- **Financial literacy gap**: SEBI's own surveys show that over 60% of retail investors cannot correctly explain the difference between a large-cap and a multi-cap fund

The Supply Side (What Exists Today)

Tool	What It Does	What It Doesn't Do
Screener.in	Fundamental data, custom screens	No portfolio view, no thesis tracking, no MF analysis
Tickertape	Stock + MF screener, MMI, portfolio tracking	No decision quality scoring, no thesis monitoring, no IPO forensics
Smallcase	Curated portfolios, thematic investing	No analysis of YOUR reasoning, no cross-asset overlap detection
Sensibull	Options analytics	Equity-only blind spot, no decision framework
Chartink	Technical screening	Pure technical, no fundamental integration
Value Research	MF analysis, overlap checker	No stock integration, no thesis tracking, basic overlap tool
Stoxra	AI-native stock analysis	New entrant, stock-only, no portfolio-level intelligence

The gap: No platform unifies stocks + mutual funds + IPOs into a single decision-quality layer. Every tool is either stock-only, MF-only, or provides data without a decision framework.

Why Now (Timing Catalysts)

1. **SEBI's 2026 mutual fund classification overhaul** — new overlap disclosure rules create demand for tools that help investors act on that data
2. **SEBI's algo trading regulations (April 2026)** — legitimized AI-assisted analysis tools within a compliance framework

3. **Monthly AMFI portfolio disclosure —** makes real-time overlap detection technically feasible with public data
4. **IPO volume surge —** more retail participation in IPOs than ever, but quality of decision-making is at an all-time low
5. **Post-2024 market correction —** investors who entered in the euphoria of 2023-24 are now sitting on losses and questioning their process for the first time

---

3. WHAT THESIS ACTUALLY IS

Product Identity

Thesis is a SaaS platform that functions as a personal investment decision layer for Indian retail investors. It sits on top of your existing brokerage and existing data tools — it doesn't replace Zerodha, Screener, or your mutual fund app. It replaces the spreadsheet, the mental notes, the Telegram screenshots, and the gut feeling.

Core Philosophy

1. **Analysis, not advisory.** Thesis never says "buy" or "sell." It says "your thesis for holding this is broken" or "your portfolio has a 42% hidden concentration in banking stocks." The user makes the decision. This is the critical SEBI compliance boundary.

2. **Process over prediction.** Thesis doesn't predict which stock will go up. It measures whether your decision-making *process* is improving over time — like a fitness tracker for your investment brain.

3. **Cross-asset truth.** Your 4 mutual funds and 12 stocks are not separate investments — they're one portfolio with one combined exposure. Thesis shows you the real picture.

4. **Thesis-first investing.** Every position in your portfolio should have a written thesis — why you bought, what needs to be true for you to hold, and what would make you sell. Thesis makes this effortless and monitors it automatically.

---

4. THE FIVE ENGINES (PRODUCT DEEP-DIVE)

ENGINE 1: Thesis Builder & Monitor

The Problem It Solves: Ask any retail investor "why did you buy Infosys?" and you'll get a vague answer — "IT sector is good" or "someone recommended it." When the stock drops 20%,

they have no framework to decide if they should hold, average down, or exit, because they never articulated what needed to be true for the investment to work.

How It Works:

Step 1 — Thesis Articulation (Guided) When a user adds a position (or imports their existing portfolio), Thesis prompts them through a structured thesis builder:

Stock: Infosys Ltd (INFY)

YOUR THESIS:

"Infosys will benefit from enterprise AI adoption driving

consulting revenue growth of 12-15% CAGR over 3 years."

KEY ASSUMPTIONS:

☑ IT services spending growth stays above 8% globally

☑ Infosys wins large deal TCV above $3B/quarter

☑ EBIT margin stays in 20-23% band

☑ USD/INR doesn't strengthen below 82

THESIS BREAKERS (auto-monitored):

⚠ If quarterly revenue growth drops below 5% for 2 quarters

⚠ If large deal TCV falls below $2B for 2 quarters

⚠ If attrition rises above 25% (talent retention risk)

⚠ If PE exceeds 30x (valuation ceiling for this thesis)

TIME HORIZON: 24 months

CONVICTION LEVEL: High (user-set)

Step 2 — Automated Monitoring Once a thesis is filed, the platform continuously monitors the key assumptions against live data:

- Quarterly results are auto-parsed from BSE/NSE filings
- Key metrics (revenue growth, margins, deal TCV, attrition) are extracted and compared against the user's stated thresholds
- If any thesis breaker triggers, the user gets a structured alert — not "Infosys fell 3%" (which every app sends) but "Your thesis assumption #2 has broken: large deal TCV was $1.8B this quarter vs your $2B threshold. This is the first quarter below your threshold."

Step 3 — Thesis Health Score Each position gets a Thesis Health Score (0-100):

- 80-100: All assumptions intact, thesis is healthy
- 50-79: One or more assumptions under stress, worth reviewing
- Below 50: Multiple thesis breakers triggered — the original reason for buying may no longer be valid

The Insight: This reframes the "should I sell?" question from emotional ("it's down 15%, I'm scared") to analytical ("my thesis is 40% broken, and here's specifically what changed").

ENGINE 2: Portfolio X-Ray (Cross-Asset Exposure)

The Problem It Solves: A typical Indian retail investor's actual holdings:

- 3-5 mutual funds (SIPs running on autopilot)
- 5-15 direct stocks
- Maybe 1-2 IPO allotments per year
- Maybe some SGBs or gold ETFs

They think they have 20+ positions. In reality, they might have effective exposure to just 8-10 companies because their MFs and direct stocks overlap massively.

How It Works:

Unified Portfolio Import:

- Connect broker accounts via Kite Connect API, Upstox API, Angel One API, Groww API (where available)
- Import mutual fund holdings via CAMS/KFintech CAS (Consolidated Account Statement) upload

- Manual entry fallback for any position
- IPO allotments from recent applications

The X-Ray Analysis:

YOUR PORTFOLIO X-RAY

Total Capital Deployed: ₹18,50,000

ACTUAL COMPANY EXPOSURE (Top 10):

1. HDFC Bank — 14.2% ⚠ HIGH

   ├─ Direct holding: 4.1%

   ├─ SBI Bluechip Fund: 3.8%

   ├─ Axis Flexi Cap: 3.2%

   └─ ICICI Pru Value Discovery: 3.1%

2. Reliance — 11.8% ⚠ HIGH

   ├─ Direct holding: 5.2%

   ├─ Nippon Large Cap: 3.4%

   └─ SBI Bluechip Fund: 3.2%

3. Infosys — 9.4% ⚠ MODERATE

   └─ ...

SECTOR CONCENTRATION:

Banking & Financial: 31.4% ⚠ SEVERE

IT Services: 18.2% ⚠ HIGH

Oil & Gas: 8.1%

Consumer: 7.3%

...

HIDDEN RISKS DETECTED:

⚠ 62% of your MF portfolio overlaps with your direct stocks

⚠ 3 of your 4 MFs have 70%+ overlap with each other

⚠ Your "diversified" portfolio behaves like a

concentrated banking+IT bet

⚠ Zero international exposure

⚠ Zero debt/fixed income allocation

Factor Exposure Analysis: Beyond just stock/sector overlap, show the user their exposure to investing factors:

- **Momentum exposure**: Are you overweight stocks that have run up 50%+ in 12 months?
- **Value exposure**: What's your portfolio-weighted PE, PB, dividend yield?
- **Quality exposure**: Portfolio-weighted ROCE, debt/equity, promoter pledge %
- **Size exposure**: Large-cap vs mid-cap vs small-cap actual split (not what the MF category says, but what the holdings actually are)

Correlation Matrix: Show how the user's holdings correlate with each other. "Your 4 mutual funds have a 0.92 average pairwise correlation — they move almost identically. You're paying 4x expense ratios for essentially one fund."

ENGINE 3: IPO Forensics

The Problem It Solves: Indian retail IPO investing in 2026 is dominated by:

- GMP (Grey Market Premium) tracking on Telegram/WhatsApp
- "Apply karo, listing gains milega" culture
- Zero analysis of DRHP (Draft Red Herring Prospectus)
- No comparison with already-listed peers
- No understanding of what the company actually does or whether the valuation makes sense

How It Works:

For every upcoming IPO, Thesis generates an automated forensic report:

Section 1 — Business Model Clarity Score

- What does this company actually do? (Plain-language summary extracted from DRHP)
- Revenue model breakdown (subscription vs one-time, B2B vs B2C, domestic vs export)
- Customer concentration risk (what % of revenue comes from top 5 clients?)
- How long has this business existed vs when did it become profitable?

Section 2 — Financial Forensics

- Revenue growth trajectory (is growth accelerating or decelerating heading into IPO?)
- Margin analysis (are margins expanding, or were they artificially inflated for the IPO window?)
- Cash flow vs profit divergence (is the company profitable on paper but burning cash?)
- Debt structure and use of IPO proceeds (fresh issue vs OFS — is the promoter cashing out?)
- Related-party transaction scan (red flag detection from DRHP disclosures)

Section 3 — Peer Comparison Grid

- Map the IPO company against 3-5 closest listed peers
- Compare: PE, PB, EV/EBITDA, revenue growth, margin profile, ROCE
- "At the IPO price band of ₹440-462, this company is valued at 38x PE. Its closest listed peer [Company X] trades at 24x PE with faster revenue growth. The IPO is priced at a 58% premium to the peer set."

Section 4 — Promoter & Governance Check

- Promoter background (prior ventures, regulatory actions, pledged shares)
- SEBI/MCA compliance history
- Auditor quality and qualifications in audit reports
- Lock-in periods and post-IPO sell pressure timeline

Section 5 — Your Portfolio Impact

- "If you get full allotment and the stock lists at GMP-implied price, it will be 2.3% of your portfolio"
- "This IPO is in the same sector as 3 stocks you already hold — your IT services exposure would go from 18% to 21%"

- Overlap analysis with existing portfolio

What Thesis Does NOT Do:

- Does NOT say "apply" or "avoid" — shows the data and lets the user decide
- Does NOT predict listing gains
- Does NOT show GMP (that's speculative noise, not analysis)

ENGINE 4: Exit Intelligence

The Problem It Solves: The hardest decision in investing is selling. Indian retail investors:

- Hold losers too long ("it'll come back") — disposition effect
- Sell winners too early ("book profits") — cutting flowers, watering weeds
- Have no systematic exit framework
- React emotionally to daily price movements

How It Works:

Pre-Commitment Exit Rules: When the user creates a thesis, they also set exit conditions — BEFORE the emotions kick in:

EXIT RULES FOR: Infosys (INFY)

THESIS-BASED EXITS:

☐ If quarterly revenue growth < 5% for 2 consecutive quarters

☐ If EBIT margin drops below 18%

☐ If large deal TCV < $2B for 2 consecutive quarters

VALUATION-BASED EXITS:

☐ If trailing PE exceeds 32x (historical 90th percentile)

☐ If stock reaches ₹2,100 (target based on thesis)

RISK-BASED EXITS:

☐ If position becomes >8% of portfolio (concentration limit)

☐ If stock drawdown exceeds 25% from purchase price (stop-loss)

PORTFOLIO-LEVEL RULES:

☐ If IT sector exposure exceeds 25% of total portfolio

☐ If portfolio correlation with Nifty IT index exceeds 0.85

Automated Monitoring & Alerts: When any exit condition approaches or triggers:

🔴 EXIT CONDITION TRIGGERED — Infosys (INFY)

RULE: "Quarterly revenue growth < 5% for 2 consecutive quarters"

STATUS: Q1 FY27 revenue growth = 3.2% (2nd consecutive quarter below 5%)

YOUR PRE-COMMITTED ACTION: Review and consider exiting

CONTEXT:

- Your thesis health score has dropped from 78 to 41

- 2 of your 3 key assumptions are now broken

- You've held for 14 months of your 24-month horizon

- Current P&L on this position: +8.2%

DECISION LOG:

What did you decide? [Hold] [Partial Exit] [Full Exit] [Override Rule]

Why? _____

Decision Journal & Pattern Analysis: Every decision the user makes (hold, sell, override) is logged. Over time, Thesis builds a pattern:

YOUR DECISION PATTERNS (Last 18 Months):

_____

Exit rules triggered: 14 times

You followed the rule: 6 times (43%)

You overrode the rule: 8 times (57%)

When you FOLLOWED your exit rules:

Average outcome: +2.3% (avoided further downside)

When you OVERRODE your exit rules:

Average outcome: -11.7% (the thesis was right to break)

INSIGHT: Your exit rules are smarter than your overrides.

Trusting your pre-committed rules would have saved you

₹1,43,000 over 18 months.

ENGINE 5: Decision Quality Score (The Meta-Layer)

The Problem It Solves: Investors judge themselves by returns, which is heavily influenced by luck and market conditions. A bull market makes bad decisions look good. A bear market makes good decisions look bad. What you actually control is *process quality* — and nobody measures it.

How It Works:

Thesis calculates a **Decision Quality Score (DQS)** based on:

1. Thesis Rigor (25%)

- Did the user articulate a clear thesis before buying?
- Were specific, measurable assumptions stated?
- Were exit conditions pre-committed?
- Score: 0 (no thesis, FOMO buy) → 100 (fully articulated thesis with measurable assumptions)

2. Portfolio Construction (25%)

- Diversification quality (not just count, but actual correlation-adjusted diversification)
- Position sizing discipline (no single position >10%, no single sector >30%)
- Asset allocation alignment with stated goals and time horizon
- Overlap management (low MF-stock overlap)

3. Exit Discipline (25%)

- Percentage of exit rules followed vs overridden
- Quality of override decisions (did overrides outperform or underperform?)
- Average holding period vs stated time horizon alignment
- Frequency of panic sells during market dips

4. Learning & Adaptation (25%)

- Are thesis quality scores improving over time?
- Are recurring mistake patterns (identified by the platform) decreasing?
- Is the user engaging with IPO forensics before applying, or skipping it?
- Is the user reviewing triggered alerts, or ignoring them?

The DQS Dashboard:

YOUR DECISION QUALITY SCORE: 67/100

	You	Avg User	Top 10%
Thesis Rigor:	72	45	88
Portfolio Const:	81	52	91
Exit Discipline:	48	38	82
Learning:	67	41	85

TREND: ↑ Improving (was 54 three months ago)

YOUR BIGGEST LEAKS:

1. Exit discipline — you override exit rules 57% of the time,

and overrides cost you money 73% of the time

2. IPO process — you applied to 6 IPOs without reviewing the

forensic report for any of them

3. Thesis gaps — 4 of your 12 stock positions have no thesis.

You're flying blind on 33% of your equity portfolio.

NEXT ACTIONS TO IMPROVE YOUR DQS:

→ Write a thesis for your 4 unthesized positions

→ Review your IT sector concentration (31% — above your 25% limit)

→ Set exit rules for your 3 positions that have none

Key Insight: This score is *independent of returns*. You can have a DQS of 90 and lose money in a bear market — that's fine, your process was sound. You can have a DQS of 30 and make money in a bull market — that's luck, not skill, and it won't last.

---

5. CROSS-ASSET UNIFIED VIEW

The killer feature is that all five engines work across ALL asset types simultaneously:

Stocks (Direct Equity)

- Full thesis tracking, exit monitoring, fundamental data
- Source: BSE/NSE APIs, quarterly results auto-parsed

Mutual Funds

- Holdings-level look-through (using AMFI monthly disclosures)
- Overlap detection with other MFs AND with direct stocks
- Expense ratio efficiency analysis ("you're paying 1.8% for a fund that's 85% identical to a 0.15% index fund")
- SIP optimization: "Your 4 SIPs collectively create a portfolio that's 73% correlated with Nifty 50. You could replace 3 of them with one index fund and one genuinely differentiated fund."

IPOs

- Pre-application forensic analysis
- Post-listing thesis tracking (did the IPO thesis play out?)
- Portfolio impact simulation before applying
- Historical accuracy tracking ("of the 8 IPOs you applied to, 5 are below listing price after 6 months")

ETFs & Index Funds

- Factor exposure analysis
- Tracking error monitoring
- Overlap with direct holdings

Government Securities & Bonds (SGBs, G-Secs via RBI Retail Direct)

- Asset allocation view — what percentage is in equity vs debt vs gold
- Rebalancing alerts when allocation drifts from target

---

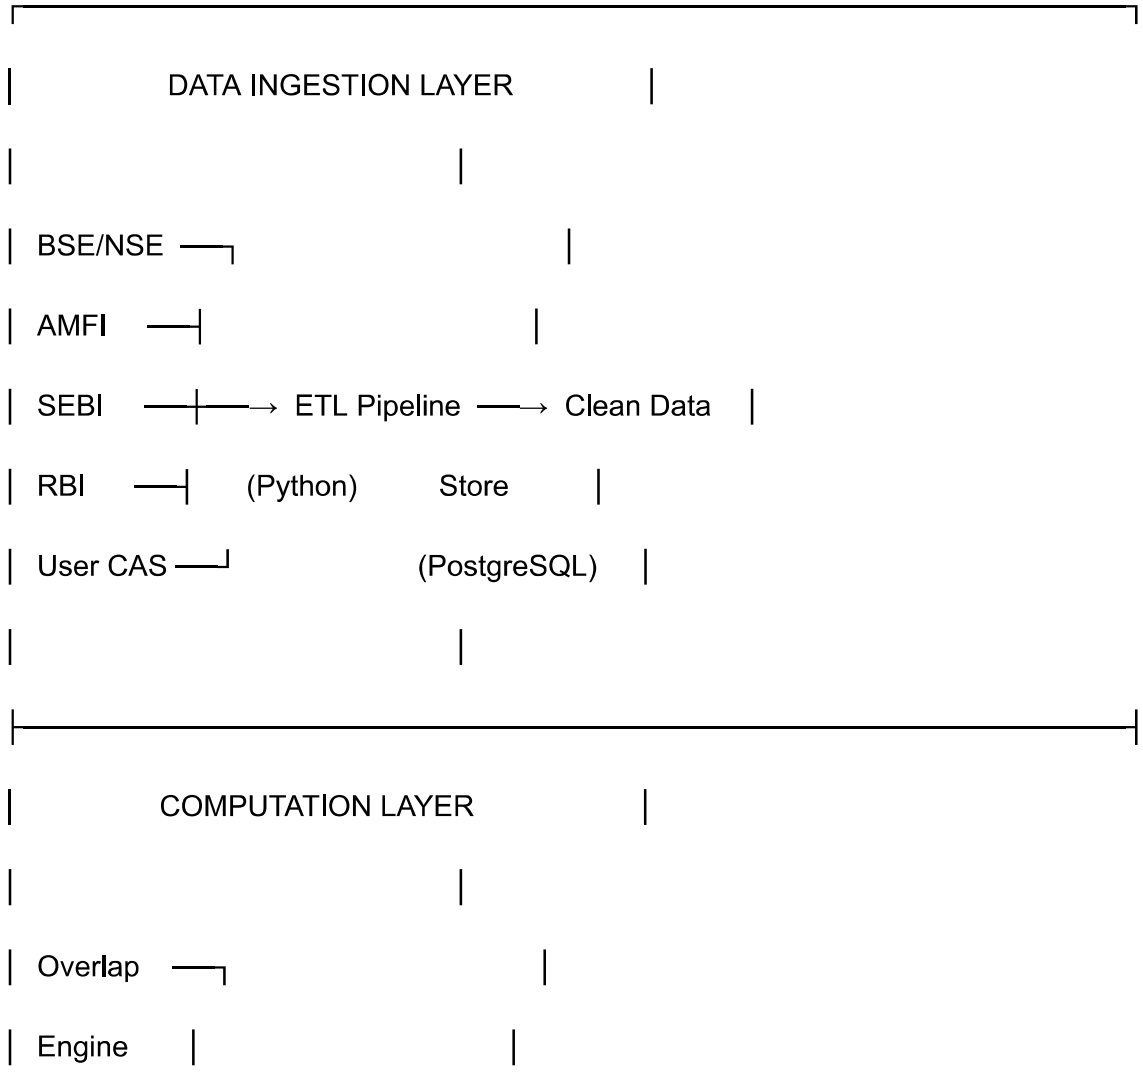
6. DATA ARCHITECTURE & SOURCES

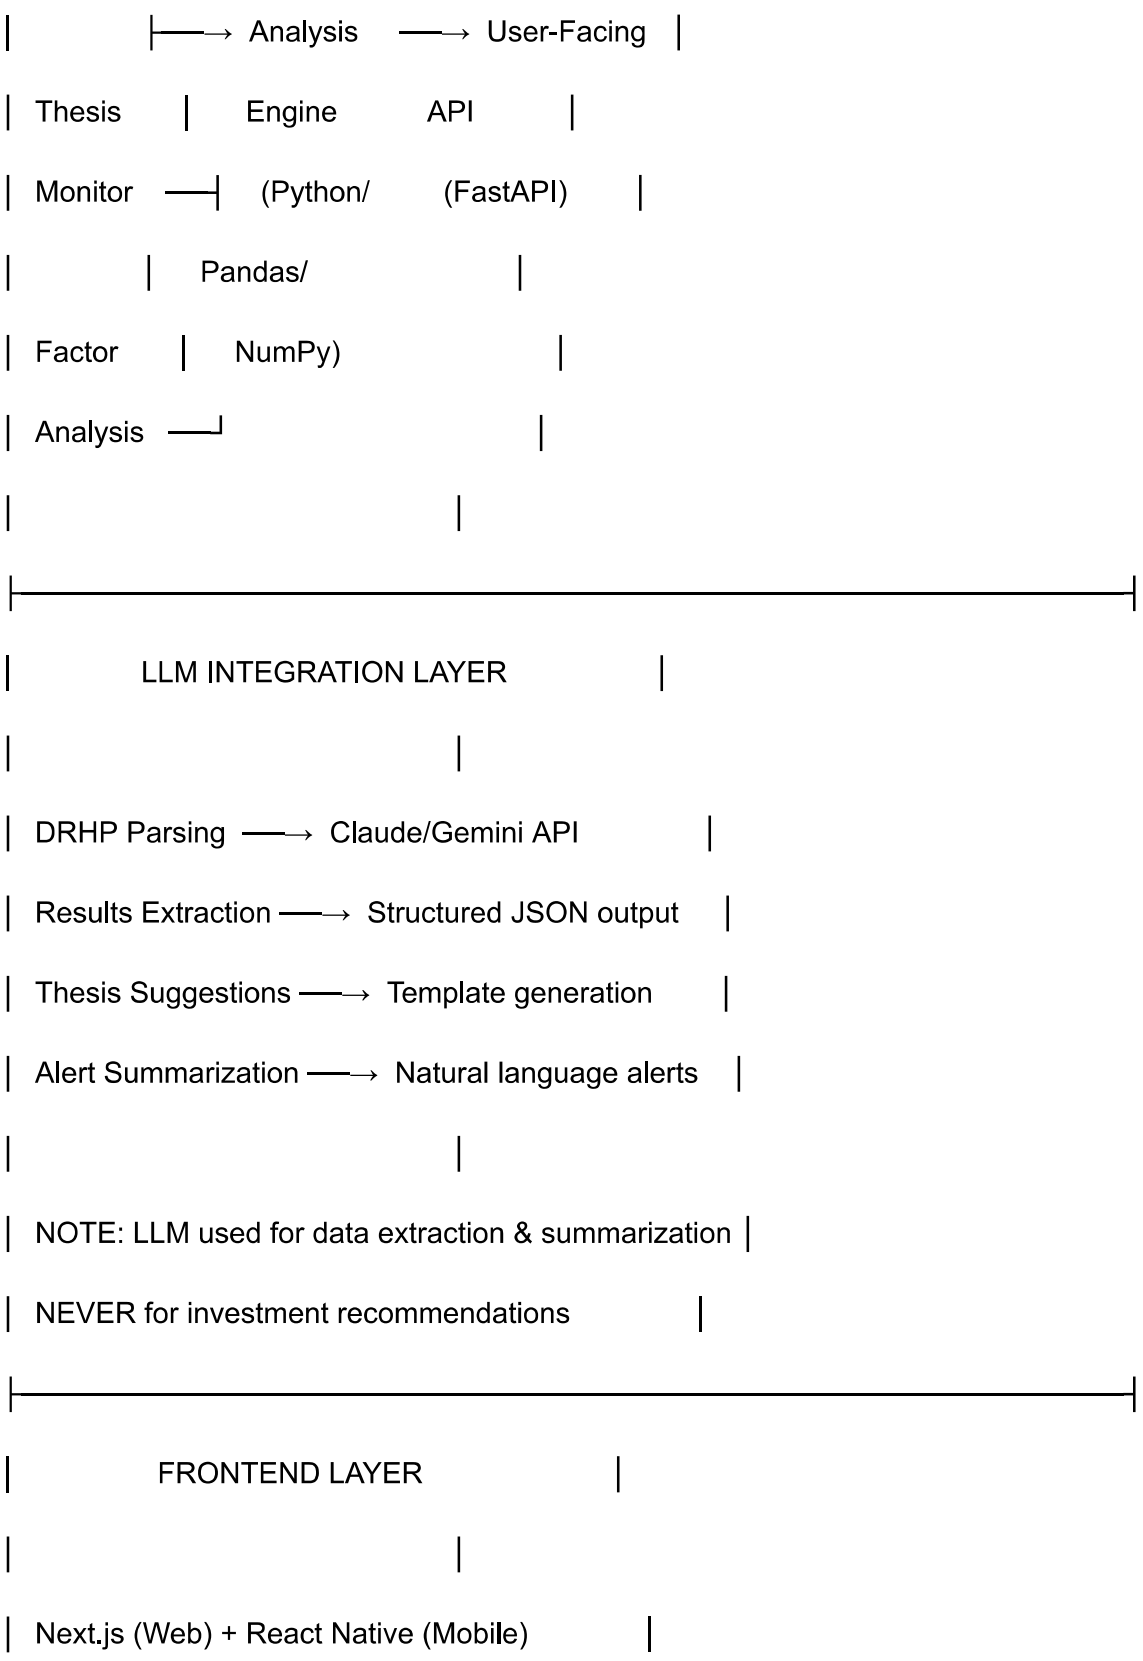
Publicly Available Data Sources (No licensing barriers)

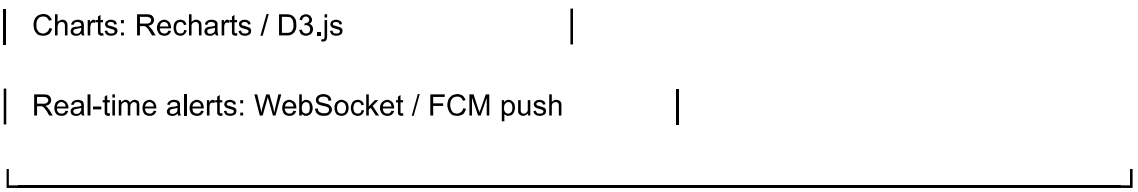
Data	**Source**	**Frequency**	**Method**
Stock prices & volumes	BSE/NSE (bhavcopy)	Daily EOD	Free download/API
Quarterly results	BSE filing portal	Quarterly	Structured scraping
MF portfolio holdings	AMFI India	Monthly	Free download
MF NAV data	AMFI India	Daily	Free API
IPO filings (DRHP/RHP)	SEBI/BSE	Per IPO	PDF parsing + LLM extraction
Corporate actions	BSE/NSE	As filed	API/scraping
Insider trading disclosures	BSE (SAST)	As filed	Structured scraping

Data	Source	Frequency	Method
Bulk/block deals	BSE/NSE	Daily	Free download
FII/DII flow data	NSDL/Moneycontrol	Daily	Structured scraping
RBI policy documents	RBI website	As released	PDF parsing
Index constituents & weights	NSE	Periodic	Free download
CAS (user portfolio)	CAMS/KFintech	User-uploaded	PDF parsing
Economic indicators	RBI DBIE, MOSPI	Monthly/Quarterly	API

Data Pipeline Architecture




Analysis
User-Facing
Thesis
Engine
API
Monitor
(Python/
(FastAPI)
Pandas/
Factor
NumPy)
Analysis
LLM INTEGRATION LAYER
DRHP Parsing
Claude/Gemini API
Results Extraction
Structured JSON output
Thesis Suggestions
Template generation
Alert Summarization
Natural language alerts
NOTE: LLM used for data extraction & summarization
NEVER for investment recommendations
FRONTEND LAYER
Next.js (Web) + React Native (Mobile)



LLM Usage (Carefully Scoped)

LLMs are used for:

- **DRHP parsing**: Extracting structured data from 400-page IPO documents (business model, financials, risk factors, related-party transactions)
- **Quarterly results extraction**: Parsing BSE filing PDFs into structured financial data
- **Natural language alert generation**: Converting "thesis breaker #2 triggered" into a human-readable explanation
- **Thesis template suggestions**: "Based on the data for [Company], here are common thesis frameworks other investors have used" (generic templates, not specific recommendations)

LLMs are NEVER used for:

- Generating buy/sell/hold recommendations
- Predicting price targets
- Telling users what to invest in
- Making personalized investment advice

---

7. SEBI COMPLIANCE MODEL

Why Thesis Doesn't Need an RA/RIA License

Under SEBI's Research Analyst Regulations (2014) and Investment Adviser Regulations (2013), the key boundary is:

Requires registration:

- Making specific buy/sell/hold recommendations on securities
- Providing personalized investment advice based on a client's financial situation
- Publishing research reports with price targets and recommendations

Does NOT require registration:

- Providing analytical tools (screeners, calculators, portfolio trackers)
- Displaying publicly available financial data in organized formats
- Offering portfolio analytics (overlap detection, diversification metrics)
- Decision journals and process tracking tools

Thesis operates in the second category. The platform:

- Never says "buy X" or "sell Y"
- Shows the user their own data (portfolio composition, overlap, thesis status)
- Provides analysis frameworks, not recommendations
- All "thesis breaker" alerts are based on rules THE USER set, not rules Thesis generated

Precedent: Screener.in, Tickertape, Chartink, Value Research all operate without RA registration because they provide tools, not advice. Thesis follows the same model with additional safeguards:

1. Every screen includes a clear disclaimer
2. The platform never generates specific security recommendations
3. IPO forensic reports show data and peer comparisons without "apply/avoid" verdicts
4. DQS scoring measures process quality, not portfolio performance

Compliance Safeguards in Product Design

- **No "top picks" or "recommendations" anywhere in the product**
- **No price targets generated by the platform —** only user-set targets tracked
- **No model portfolios or suggested allocations —** only analysis of the user's existing portfolio
- **Mandatory risk disclaimers** on every analytical output
- **Audit trail** of every piece of analysis shown to users, in case of regulatory inquiry

---

8. MONETIZATION & UNIT ECONOMICS

Pricing Tiers

Free Tier — "Lens"

- Track up to 5 direct stock positions with basic thesis builder
- Portfolio X-Ray for up to 2 mutual funds
- 1 IPO forensic report per month
- Basic DQS score
- Purpose: top-of-funnel acquisition, habit formation

Pro Tier — ₹349/month (₹2,999/year)

- Unlimited stock thesis tracking
- Full portfolio X-Ray across all asset types
- Unlimited IPO forensic reports
- Complete exit intelligence engine
- Full DQS with historical trends and pattern analysis
- Email + push alerts for thesis breakers
- Target user: Serious retail investor with ₹5-50L portfolio

Premium Tier — ₹799/month (₹6,999/year)

- Everything in Pro
- Multi-portfolio tracking (personal + spouse + HUF)
- Advanced factor exposure analysis
- Backtesting: "If I had followed my exit rules consistently for the last 3 years, what would my returns have been?"
- CSV/PDF export of all reports
- API access for power users
- Target user: HNI retail or semi-professional investor with ₹50L+ portfolio

B2B Tier — ₹5,000-15,000/month per advisor seat

- White-labeled version for SEBI-registered RIAs and wealth managers
- Client portfolio X-Ray across all clients
- Thesis tracking across client portfolios
- Generate client-facing reports: "Here's why we recommended this, here are the thesis conditions we're monitoring, here's the current health score"
- Compliance documentation: auto-generated audit trail of every recommendation rationale
- Target user: Small and mid-sized RIA firms (500+ registered RIAs in India)

Unit Economics (Projected)

Metric	**Year 1**	**Year 2**	**Year 3**
Free users	50,000	200,000	600,000
Paid users	2,500	15,000	60,000
Conversion rate	5%	7.5%	10%

Metric	Year 1	Year 2	Year 3
ARPU (blended)	₹300/mo	₹350/mo	₹400/mo
Monthly revenue	₹7.5L	₹52.5L	₹2.4Cr
B2B seats	20	150	500
B2B revenue	₹1.5L/mo	₹12L/mo	₹50L/mo
Gross margin	75%	80%	82%

Key cost drivers: LLM API costs (DRHP parsing), cloud infrastructure, data pipeline maintenance. LLM costs are per-analysis (not per-user), so they scale sublinearly with user growth.

---

9. COMPETITIVE LANDSCAPE & MOAT

Direct Competitors: None

No platform in India combines thesis tracking + cross-asset X-ray + IPO forensics + exit intelligence + decision quality scoring. These features exist in isolation:

- MF overlap: Value Research (basic, 2-fund comparison only)
- Stock screening: Screener.in, Tickertape (data, not decision frameworks)
- IPO info: Chittorgarh IPO, IPO Watch (listing data, not forensic analysis)
- Portfolio tracking: INDmoney, Groww (tracking, not analysis)
- Decision journaling: Nothing in India (Composer in the US is closest, but US-only and algo-focused)

The Moat (Defensibility Over Time)

1. Data Network Effect (Strongest Moat) Every user who tracks theses, follows/overrides exit rules, and logs decisions contributes to the anonymized pattern database. Over time, Thesis can show:

- "Investors who override exit rules in IT stocks lose money 73% of the time"
- "IPOs where promoter OFS exceeds 60% of issue size underperform peers 80% of the time within 12 months"
- "The most common thesis for [Stock X] among Thesis users has a 65% thesis health rate"

This behavioral data is unique and impossible to replicate without the user base.

2. User Lock-In (Decision History) A user's 2-year decision journal — their theses, exit rules, overrides, patterns, DQS history — is deeply personal and non-portable. Switching costs increase with every logged decision.

3. LLM-Powered DRHP Parsing (Technical Moat) Building a reliable pipeline to extract structured data from 400-page legal-financial documents (DRHPs) is genuinely hard. Hallucination in financial data extraction is a critical accuracy problem. Getting this right requires significant prompt engineering, validation layers, and domain-specific fine-tuning.

4. Cross-Asset Integration (Complexity Moat) Unifying BSE/NSE stock data + AMFI mutual fund holdings data + SEBI IPO filings + user portfolio data from multiple brokers into a single coherent view is an integration challenge that takes 12-18 months to get right. Each data source has its own format, update frequency, and quirks.

---

10. GO-TO-MARKET STRATEGY

Phase 1: Content-Led Growth (Months 0-6)

Twitter/X Fintwit Community: Indian fintwit is massive and engaged. Strategy:

- Publish weekly "Portfolio X-Ray" case studies (anonymized user portfolios showing hidden concentrations)
- Publish "IPO Forensic" reports for every upcoming IPO — position Thesis as the place for serious IPO analysis
- Run a "Decision Quality Challenge" — users share their investment mistakes and the Thesis framework for avoiding them

YouTube:

- "I X-Rayed My Portfolio and Found This" — visual walkthroughs of overlap detection
- "The IPO Nobody Read the DRHP For" — forensic breakdowns of hyped IPOs
- Partner with Indian finance YouTubers (Akshat Shrivastava, Pranjal Kamra, CA Rachana Ranade) for exposure

Why content works here: Indian retail investors are hungry for structured frameworks. The current content landscape is dominated by "top 5 stocks to buy" clickbait. Rigorous,

framework-driven content stands out immediately and builds trust with the exact audience that will pay for the product.

Phase 2: Community & Virality (Months 6-12)

Shareable Analytics:

- "Share your Portfolio X-Ray" — a visual summary card (like Spotify Wrapped for your portfolio) showing your DQS, top concentration risks, and thesis health
- "My IPO Track Record" — how many IPOs you applied to vs how many are profitable 6 months later
- These cards are designed for Twitter/WhatsApp sharing — organic viral loop

Referral Program:

- "Invite 3 friends → get 1 month Pro free" — standard but effective for a tool where word-of-mouth matters

Phase 3: B2B Expansion (Months 12-18)

Target: SEBI-Registered RIAs (Research Analysts & Investment Advisors)

- India has 1,000+ registered RIAs, most running small practices (1-50 clients)
- Their biggest pain: documenting WHY they recommend something and tracking the recommendation over time
- Thesis B2B gives them: client portfolio X-Ray, thesis documentation per recommendation, exit monitoring, compliance-ready audit trail
- Sales motion: direct outreach + conference presence (CFA Society India, ARIA events)

---

11. TECHNICAL ARCHITECTURE

Stack

Layer	Technology	Why
Frontend (Web)	Next.js 14 + TypeScript	SSR for SEO, React ecosystem for interactive charts

Layer	Technology	Why
Frontend (Mobile)	React Native / Expo	Code sharing with web, critical for Indian users (80%+ mobile)
API Layer	Python + FastAPI	Fast development, async support, great for data-heavy applications
Database	PostgreSQL + TimescaleDB extension	Relational for portfolio data, timeseries for price/NAV history
Cache	Redis	Session management, frequent query caching
Task Queue	Celery + Redis	Async processing of DRHP parsing, portfolio recalculation
LLM Integration	Claude API (primary) + Gemini (fallback)	DRHP parsing, results extraction, alert generation
Search	Meilisearch	Fast stock/MF/IPO search across the platform
Infrastructure	AWS Mumbai region (ap-south-1)	Data residency in India, low latency for Indian users
Monitoring	PostHog + Sentry	Product analytics + error tracking
CI/CD	GitHub Actions	Standard

Key Technical Challenges & Solutions

Challenge 1: DRHP Parsing Accuracy

- DRHPs are 300-600 page PDFs with complex table structures
- Solution: Multi-pass extraction — first pass with LLM for structured data, second pass with rule-based validation (e.g., verify that extracted revenue figures match the sum of segment revenues), human review for first 50 IPOs to build confidence
- Accuracy target: 99%+ on financial figures (zero tolerance for wrong numbers in financial context)

Challenge 2: MF Holdings Reconciliation

- AMFI publishes MF holdings monthly, but format varies by AMC
- Some AMCs publish partial holdings (only top 10)
- Solution: Standardized parser per AMC, fallback to category-level exposure estimation when full holdings unavailable

Challenge 3: Real-Time Thesis Monitoring

- Users may have hundreds of thesis conditions to monitor
- Solution: Batch processing — run thesis checks once daily after market close (most thesis conditions are based on quarterly results or daily price, not intraday), with priority queue for conditions near threshold

Challenge 4: Privacy & Data Security

- Portfolio data is highly sensitive financial information
- Solution: End-to-end encryption for portfolio data, SOC2 Type 2 certification by Year 2, data stored in India (AWS Mumbai), clear privacy policy with no third-party data sharing

---

12. MVP SCOPE & PHASED ROADMAP

MVP (Months 1-4): "Thesis for Stocks"

Build:

- Thesis Builder (guided thesis creation for direct stocks)
- Basic thesis health monitoring (manual quarterly results input + BSE auto-parsing for top 200 stocks)
- Simple exit rule setting and alerting
- Basic DQS score (thesis rigor component only)
- Web app only

Don't build (yet):

- Mutual fund integration
- IPO forensics
- Cross-asset X-Ray
- Mobile app

- B2B features

Target: 500-1,000 beta users from fintwit, validate that thesis tracking is a behavior users will adopt.

Key validation question: Will users actually write theses for their holdings? If yes → everything else follows. If no → the product needs a fundamentally different entry point.

Phase 2 (Months 4-8): "Portfolio X-Ray"

Add:

- CAS upload for mutual fund import
- Cross-asset overlap detection (MF + stocks)
- Sector concentration analysis
- Factor exposure analysis
- Basic portfolio-level DQS
- Start building mobile app

Key validation question: Does showing users their real exposure change their behavior? Do they actually sell redundant MFs or reduce concentration?

Phase 3 (Months 8-12): "IPO Forensics + Exit Intelligence"

Add:

- Automated DRHP parsing and forensic report generation
- Full exit intelligence engine with decision logging
- Decision pattern analysis ("your overrides cost you ₹X")
- Complete DQS with all 4 components
- Mobile app launch
- Broker API integrations (Kite Connect first, then Upstox, Groww)

Phase 4 (Months 12-18): "B2B + Scale"

Add:

- B2B RIA platform (white-labeled)
- Multi-portfolio support
- Backtesting engine
- API for power users
- Community features (anonymized thesis sharing, DQS leaderboards)

---

13. RISK REGISTER

Risk	Severity	Probability	Mitigation
SEBI regulatory action (classified as advisory)	Critical	Low	Conservative product design, legal review, no recommendations ever
Users don't write theses (behavior change fails)	High	Medium	Make thesis creation extremely easy (template-driven, 2-minute process), show immediate value (thesis health score)
Data source disruption (BSE/NSE changes format/access)	Medium	Medium	Multiple data sources for redundancy, maintain parsers actively
LLM accuracy on DRHP parsing	High	Medium	Multi-pass validation, human review layer, accuracy monitoring
Incumbent launches similar feature	Medium	High	Speed of execution, behavioral data moat (they have data, you have decision history)
Low conversion from free to paid	High	Medium	Gate high-value features (exit intelligence, full X-Ray) behind paywall, ensure free tier creates clear "I need more" moments
Privacy breach of portfolio data	Critical	Low	Encryption, SOC2, security audits, minimal data retention policy

---

14. TEAM COMPOSITION (Initial: 4-6 People)

Role	Focus	Skills
Founder/CEO	Product vision, GTM, fundraising	Finance domain knowledge + product sense
Full-Stack Engineer #1	Backend + data pipeline	Python, FastAPI, PostgreSQL, data engineering
Full-Stack Engineer #2	Frontend + mobile	Next.js, React Native, D3.js/Recharts
Data/LLM Engineer	DRHP parsing, data extraction, NLP pipeline	LLM integration, prompt engineering, PDF parsing
(Phase 2) Designer	Product design, shareable cards, mobile UX	UI/UX with fintech experience
(Phase 2) Growth	Content, community, social	Fintwit presence, finance content creation

---

15. APPENDIX: KEY INDIAN MARKET DATA SOURCES

Free / Public Sources

Source	URL	Data Available
BSE India	bseindia.com	Price data, corporate filings, insider trading, bulk deals
NSE India	nseindia.com	Price data, index data, FII/DII activity
AMFI India	amfiindia.com	MF NAV, portfolio holdings (monthly), scheme details

Source	URL	Data Available
SEBI	sebi.gov.in	IPO filings, regulatory orders, RA/RIA registry
MCA	mca.gov.in	Company filings, director info, charge registrations
RBI DBIE	dbie.rbi.org.in	Macro data, interest rates, forex data
Trendlyne	trendlyne.com	Financial data API (freemium)
CAMS/KFintech	camsonline.com / kfintech.com	CAS statements (user-downloaded)

Paid Data Sources (For Scale)

Source	Data	Approximate Cost
BSE StAR MF API	Real-time MF data	₹50K-2L/year
NSE Data Feed	Real-time price feed	₹1-5L/year
CMIE Prowess	Comprehensive financials	₹3-8L/year
Capitaline	Historical financials	₹2-5L/year

SUMMARY: WHY THESIS WINS

1. **Novel positioning**: Not "what to buy" but "is your reasoning good?" — no competitor occupies this space
2. **Regulatory moat**: By design, Thesis cannot be classified as advisory — it's a decision-quality tool
3. **Cross-asset truth**: First platform to unify stocks + MFs + IPOs into one exposure view with behavioral analytics
4. **Behavioral data moat**: Decision journals, thesis patterns, and exit-rule compliance data create a unique dataset that grows more valuable with every user
5. **India-first, India-specific**: Built for SEBI regulations, NSE/BSE data, AMFI disclosures, and Indian investor behavior patterns — not a localized global product
6. **Clear monetization path**: Consumer SaaS (₹349-799/month) + B2B to RIAs (₹5K-15K/month) — two revenue engines from the same core product

---

This document is a product ideation framework. It does not constitute investment advice, financial guidance, or a solicitation to invest. All references to financial instruments, regulations, and market data are for product design purposes only.